

OMG -

let's build a **shell** in 10 mins

Stefanie Schirmer

@linse

OMG -

let's build a **shell** in 10 mins

Stefanie Schirmer

@linse

What is a shell?

```
#include <stdio.h> // printf
#include <string.h> // strlen
#include <stdlib.h> // exit
#include <unistd.h> // syscalls

int MAXLINE = 1024;
char prompt[] = "lnsh> "; // command line prompt

int main(int argc, char **argv) {
    char cmdline[MAXLINE]; // buffer for fgets

    while (1) {

        printf("%s", prompt);

        if ((fgets(cmdline, MAXLINE, stdin) == NULL) && ferror(stdin))
            error("fgets error");

        if (feof(stdin)) {
            printf("\n");
            exit(0);
        }

        // remove trailing newline
        cmdline[strlen(cmdline)-1] = '\0';

        // evaluate command line
        eval(cmdline);
    }
}
```

show prompt + read input

call parse

```
void eval(char *cmdline) {
    int bg;
    struct command cmd;

    printf("Evaluating '%s'\n", cmdline);

    // parse cmdline into cmd structure
    bg = parse(cmdline, &cmd);
    printf("Found command %s\n", cmd.argv[0]);
    printf("Builtin %d\n", cmd.builtin);

    // -1 means parse error
    if (bg == -1) return;
}
```

command struct

```
// maxargs in struct can't be variable, use define
#define MAXARGS 128

struct command {
    int argc;           // number of args
    char *argv[MAXARGS]; // arguments list
    enum builtin_t {    // is argv[0] a builtin command?
        NONE, QUIT, JOBS, BG, FG } builtin;
};
```

```

int parse(const char *cmdline, struct command *cmd) {
    static char array[MAXLINE];           // local copy of command line
    const char delims[10] = " \t\r\n";   // argument delimiters
    char *buf = array;                    // ptr that traverses command line
    char *next;                           // ptr to the end of the current arg
    char *endbuf;                         // ptr to the end of the cmdline string
    int is_bg;                            // background job?

    if (cmdline == NULL) {
        (void) fprintf(stderr, "Error: command line is NULL\n");
        return -1;
    }

    (void) strncpy(buf, cmdline, MAXLINE);
    endbuf = buf + strlen(buf);

    // build argv list
    cmd->argc = 0;

    while (buf < endbuf) {
        // skip delimiters
        buf += strspn (buf, delims);
        if (buf >= endbuf) break;

        // Find next delimiter
        next = buf + strcspn (buf, delims);

        // terminate the token
        *next = '\0';

        // Record token as the next argument
        cmd->argv[cmd->argc++] = buf;

        // Check if argv is full
        if (cmd->argc >= MAXARGS-1) break;

        buf = next + 1;
    }

    // argument list must end with a NULL pointer
    cmd->argv[cmd->argc] = NULL;

    if (cmd->argc == 0) // ignore blank line
        return 1;

    cmd->builtin = parseBuiltin(cmd);

    // should job run in background?
    if ((is_bg = (*cmd->argv[cmd->argc-1] == '&')) != 0)
        cmd->argv[--cmd->argc] = NULL;

    return is_bg;
}

```

parse():

line -> struct

builtin shell commands

```
enum builtin_t parseBuiltin(struct command *cmd) {
    if (!strcmp(cmd->argv[0], "quit")) { // quit command
        return QUIT;
    } else if (!strcmp(cmd->argv[0], "jobs")) { // jobs command
        return JOBS;
    } else if (!strcmp(cmd->argv[0], "bg")) { // bg command
        return BG;
    } else if (!strcmp(cmd->argv[0], "fg")) { // fg command
        return FG;
    } else {
        return NONE;
    }
}
```


System calls

run programs

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main (int argc, char const *argv[]) {
    printf("Hi there\n");
```

```
    // executing an ls command
    char* argv2[] = {"ls", "-la", NULL};
    execvp(argv2[0], argv2);
```

```
    printf(".. and done!");
}
```

fork + children

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
```

```
int main() {
    pid_t childPid;
    printf("Before fork %d\n", getpid());
```

```
// fork creates a child and returns its process ID
    childPid = fork();
```

```
    printf("After fork %d\n", getpid());
    exit(0);
}
```

Fork + *execvp*

she11.c

```
// updated!!!
void eval(char *cmdline) {
    int bg;
    struct command cmd;

    printf("Evaluating '%s'\n", cmdline);
    // parse line into command struct
    bg = parse(cmdline, &cmd);

    // parse error
    if (bg == -1) return;
    // empty line - ignore
    if (cmd.argv[0] == NULL) return;

    if (cmd.builtin == NONE)
        runSystemCommand(&cmd, bg);
    else
        runBuiltinCommand(&cmd, bg);
}
```

evaluate for real

fork + execvp

= shell

```
void runSystemCommand(struct command *cmd, int bg) {
    pid_t childPid;
    // FORK !!!
    if ((childPid = fork()) < 0)
        error("fork() error");
    else if (childPid == 0) { // I'm the child and could run a command
        // EXECVP !!!!
        if (execvp(cmd->argv[0], cmd->argv) < 0) {
            printf("%s: Command not found\n", cmd->argv[0]);
            exit(0);
        }
    }
    else { // I'm the parent. Shell continues here.
        if (bg)
            printf("Child in background [%d]\n", childPid);
        else
            wait(&childPid);
    }
}
```

builtin shell commands

```
void runBuiltinCommand(struct command *cmd, int bg) {  
    switch (cmd->builtin) {  
        case QUIT:  
            printf("TODO: quit\n"); break;  
        case JOBS:  
            printf("TODO: jobs\n"); break;  
        case BG:  
            printf("TODO: background\n"); break;  
        case FG:  
            printf("TODO: foreground\n"); break;  
        default:  
            error("Unknown builtin command");  
    }  
}
```

```

#include <stdio.h> // printf
#include <string.h> // strlen
#include <stdlib.h> // exit
#include <unistd.h> // syscalls

#define MAXLINE 1024
// maxargs used in struct, can't be variable so we have to use define
#define MAXARGS 128

char prompt[] = "linsh> "; // command line prompt

struct command {
    int argc; // number of args
    char *argv[MAXARGS]; // arguments list
    enum builtin_t { // is argv[0] a builtin command?
        NONE, QUIT, JOBS, BG, FG } builtin;
};

enum builtin_t parseBuiltin(struct command *cmd) {
    if (strcmp(cmd->argv[0], "quit")) { // quit command
        return QUIT;
    } else if (strcmp(cmd->argv[0], "jobs")) { // jobs command
        return JOBS;
    } else if (strcmp(cmd->argv[0], "bg")) { // bg command
        return BG;
    } else if (strcmp(cmd->argv[0], "fg")) { // fg command
        return FG;
    } else {
        return NONE;
    }
}

int parse(const char *cmdline, struct command *cmd) {
    static char array[MAXLINE]; // local copy of command line
    const char delims[10] = " \t\r\n"; // argument delimiters
    char *buf = array; // ptr that traverses command line
    char *next; // ptr to the end of the current arg
    char *endbuf; // ptr to the end of the cmdline string
    int is_bg; // background job?

    if (cmdline == NULL) {
        (void) fprintf(stderr, "Error: command line is NULL\n");
        return -1;
    }

    (void) strncpy(buf, cmdline, MAXLINE);
    endbuf = buf + strlen(buf);

    // build argv list
    cmd->argc = 0;

    while (buf < endbuf) {
        // skip delimiters
        buf += strspn(buf, delims);
        if (buf >= endbuf) break;

        // Find next delimiter
        next = buf + strcspn(buf, delims);

        // terminate the token
        *next = '\0';

        // Record token as the next argument
        cmd->argv[cmd->argc++] = buf;

        // Check if argv is full
        if (cmd->argc >= MAXARGS-1) break;

        buf = next + 1;
    }

    // argument list must end with a NULL pointer
    cmd->argv[cmd->argc] = NULL;

    if (cmd->argc == 0) // ignore blank line
        return 1;

    cmd->builtin = parseBuiltin(cmd);

    // should job run in background?
    if ((is_bg = (cmd->argv[cmd->argc-1] == '&')) != 0)
        cmd->argv[--cmd->argc] = NULL;

    return is_bg;
}

void error(char *msg) {
    printf("Error: %s", msg);
    exit(0);
}

void runSystemCommand(struct command *cmd, int bg) {
    pid_t childPid;
    // Fork !!
    if ((childPid = fork()) < 0)
        error("fork() error");
    else if (childPid == 0) { // I'm the child and could run a command
        // EXECVE !!!
        if (execvp(cmd->argv[0], cmd->argv) < 0) {
            printf("ls: Command not found\n", cmd->argv[0]);
            exit(0);
        }
    } else { // I'm the parent. Shell continues here.
        if (bg)
            printf("Child in background [%d]\n", childPid);
        else
            wait(&childPid);
    }
}

void runBuiltinCommand(struct command *cmd, int bg) {
    switch (cmd->builtin) {
        case QUIT:
            printf("TODO: quit\n"); break;
        case JOBS:
            printf("TODO: jobs\n"); break;
        case BG:
            printf("TODO: background\n"); break;
        case FG:
            printf("TODO: foreground\n"); break;
        default:
            error("Unknown builtin command");
    }
}

// We changed this
void eval(char *cmdline) {
    int bg;
    struct command cmd;

    printf("Evaluating '%s'\n", cmdline);
    // parse line into command struct
    bg = parse(cmdline, &cmd);

    // parse error
    if (bg == -1) return;
    // empty line - ignore
    if (cmd.argv[0] == NULL) return;

    if (cmd.builtin == NONE)
        runSystemCommand(&cmd, bg);
    else
        runBuiltinCommand(&cmd, bg);
}

// We already know this ..
int main(int argc, char **argv) {
    char cmdline[MAXLINE]; // buffer for fgets
    char c;

    while (1) {
        printf("%s", prompt);

        if ((fgets(cmdline, MAXLINE, stdin) == NULL) && !error(stdin))
            error("fgets error");

        if (feof(stdin)) {
            printf("\n");
            exit(0);
        }

        // Remove trailing newline
        cmdline[strlen(cmdline)-1] = '\0';

        // Evaluate command line
        eval(cmdline);
    }
}

```

done!

Thank you!

I'm Stefanie Schirmer

sschirme@gmail.com

[@linse](#) on twitter